

Phase 1: Data Preparation & EDA in Spark

CS236 — Database Management System

Pankaj Sharma (SID: 862549035)
Saransh Gupta (SID: 862548920)

Fall 2025

0. Executive Summary

In this phase, we set up a local PySpark environment, used two CSV datasets (`customer-reservations.csv` and `hotel-booking.csv`), performed exploratory data analysis (EDA), and data preprocessing.

1. Installation Process

1.1 Unpack the Project

```
unzip CS236-Project.zip  
cd CS236-Project
```

1.2 Creating a Virtual Environment

```
python3 -m venv venv  
source venv/bin/activate
```

1.3 Install Dependencies

```
pip install -r requirements.txt
```

1.4 Installing Java 11

```
brew install openjdk@11  
export JAVA_HOME=/opt/homebrew/opt/openjdk@11/libexec/openjdk.jdk/Contents/Home  
export PATH="/opt/homebrew/opt/openjdk@11/bin:$PATH"
```

1.5 Running the Script

```
export JAVA_HOME=/opt/homebrew/opt/openjdk@11/libexec/openjdk.jdk/Contents/Home \  
&& export PATH="/opt/homebrew/opt/openjdk@11/bin:$PATH" \  
&& spark-submit phase1-eda.py
```

2. Exploratory Data Analysis Process

This section contains EDA process and findings that we performed on the datasets - `customer-reservations.csv` and `hotel-booking.csv`. Here we will answer What, Why and Insights we gained from the EDA steps that we performed.

2.1 Schema Inspection

Customer Reservations Schema

Booking_ID	string	True
stays_in_weekend_nights	integer	True
stays_in_week_nights	integer	True
lead_time	integer	True
arrival_year	integer	True
arrival_month	integer	True
arrival_date	integer	True
market_segment_type	string	True
avg_price_per_room	double	True
booking_status	string	True

Hotel Bookings Schema

hotel	string	True
booking_status	integer	True
lead_time	integer	True
arrival_year	integer	True
arrival_month	string	True
arrival_date_week_number	integer	True
arrival_date_day_of_month	integer	True
stays_in_weekend_nights	integer	True
stays_in_week_nights	integer	True
market_segment_type	string	True
country	string	True
avg_price_per_room	double	True
email	string	True

Why:

We did this step to confirm proper loading of datasets using PySpark and to get brief idea about the dataset like columns and its data types.

Insights:

- *Customer Reservations:* This Dataset contains columns of different data types. We have Numeric fields such as `lead_time`, `arrival_year`, and `avg_price_per_room`, and categorical fields like `market_segment_type` and `booking_status` were `StringType`.
 - *Hotel Bookings:* This Dataset contains columns of 2 data types. `arrival_month` was inferred as a `StringType` (e.g., “July”, “September”), while other date components were integers.
-

2.2 Size Inspection

Why:

We did to check the size of the database in order to fit within Spark in memory limits.

Insights:

- *Customer Reservations:* **36,275 rows × 10 columns** — each row represents one unique booking (`Booking_ID`).
- *Hotel Bookings:* **78,703 rows × 13 columns** — includes additional contextual fields such as `hotel`, `country`, and `email`.
- Both datasets are of moderate size and fit comfortably within Spark’s in-memory analysis limits.

2.3 Missing Values

Missing Values Count — Customer Reservations

Booking_ID	0
stays_in_weekend_nights	0
stays_in_week_nights	0
lead_time	0
arrival_year	0
arrival_month	0
arrival_date	0
market_segment_type	0
avg_price_per_room	0
booking_status	0

Note: Customer Reservations Dataset does not have `arrival_date_week_number` field. However, during data preprocessing, we calculate it from the existing date components (year, month, day) using ISO 8601 week numbering to ensure compatibility with the Hotel Bookings Dataset.

Missing Values Count — Hotel Bookings

hotel	0
booking_status	0
lead_time	0
arrival_year	0
arrival_month	0
arrival_date_week_number	0
arrival_date_day_of_month	0
stays_in_weekend_nights	0
stays_in_week_nights	0
market_segment_type	0
country	405
avg_price_per_room	0
email	0

Why:

We did this step to find incomplete fields to decide whether to remove them or to fill in mean value.

Insights:

- *Customer Reservations:* **0 missing values** across all the columns.
- *Hotel Bookings:* Only the `country` column contained **405 missing values** all other columns were complete.

- Overall both the datasets were almost complete with only Hotel Bookings Dataset having very less percentage of missing values.
-

2.4 Distinct Values Inspection

Distinct Values Count — Customer Reservations

Booking_ID	36,275
stays_in_weekend_nights	8
stays_in_week_nights	18
lead_time	352
arrival_year	2
arrival_month	12
arrival_date	31
market_segment_type	5
avg_price_per_room	3,930
booking_status	2

Distinct Values Count — Hotel Bookings

hotel	2
booking_status	2
lead_time	439
arrival_year	2
arrival_month	12
arrival_date_week_number	53
arrival_date_day_of_month	31
stays_in_weekend_nights	17
stays_in_week_nights	32
market_segment_type	8
country	159
avg_price_per_room	6,985
email	77,144

Why:

This is essential EDA process which results in determining primary keys, redundancy and also helps in process of prediction.

Insights:

- Customer Reservations:*
 - Booking_ID had **36,275 unique values**, which means every row can be uniquely identified from the Booking_ID hence it can be used as a primary.

- ▶ `market_segment_type` had **5 categories** and can be used to draw significant inference from the data.
 - ▶ `booking_status` had **2 values** — *Canceled* and *Not_Canceled*.
 - ▶ No duplicate rows found.
 - *Hotel Bookings*:
 - ▶ `hotel` had 2 values (*City Hotel*, *Resort Hotel*).
 - ▶ `market_segment_type` had 8 categories.
 - ▶ `country` had 159 unique codes.
 - ▶ `email` had 77,144 unique entries, nearly matching total rows, implying one booking per customer.
-

2.5 Correlation Inspection

Correlation map - Customer Reservations

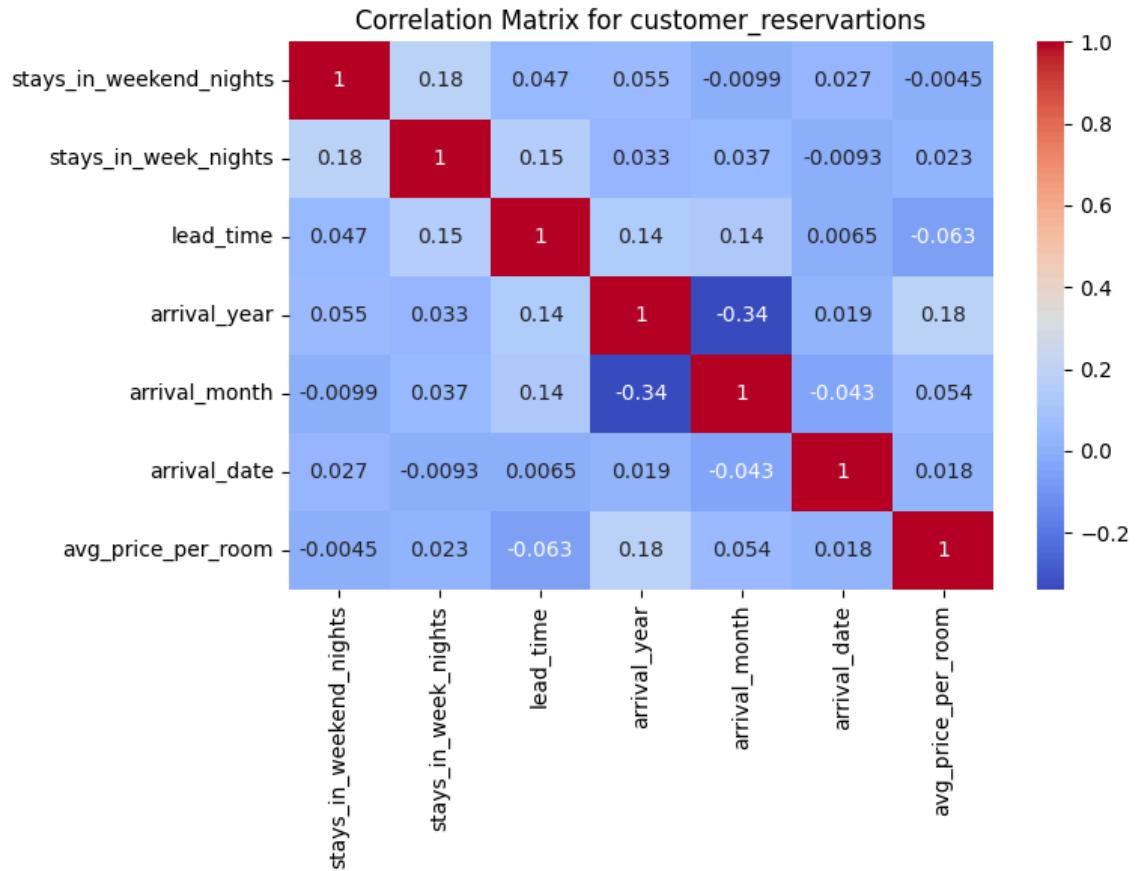


Figure 1: Correlation for Customer Reservation

Correlation map - Hotel Bookings

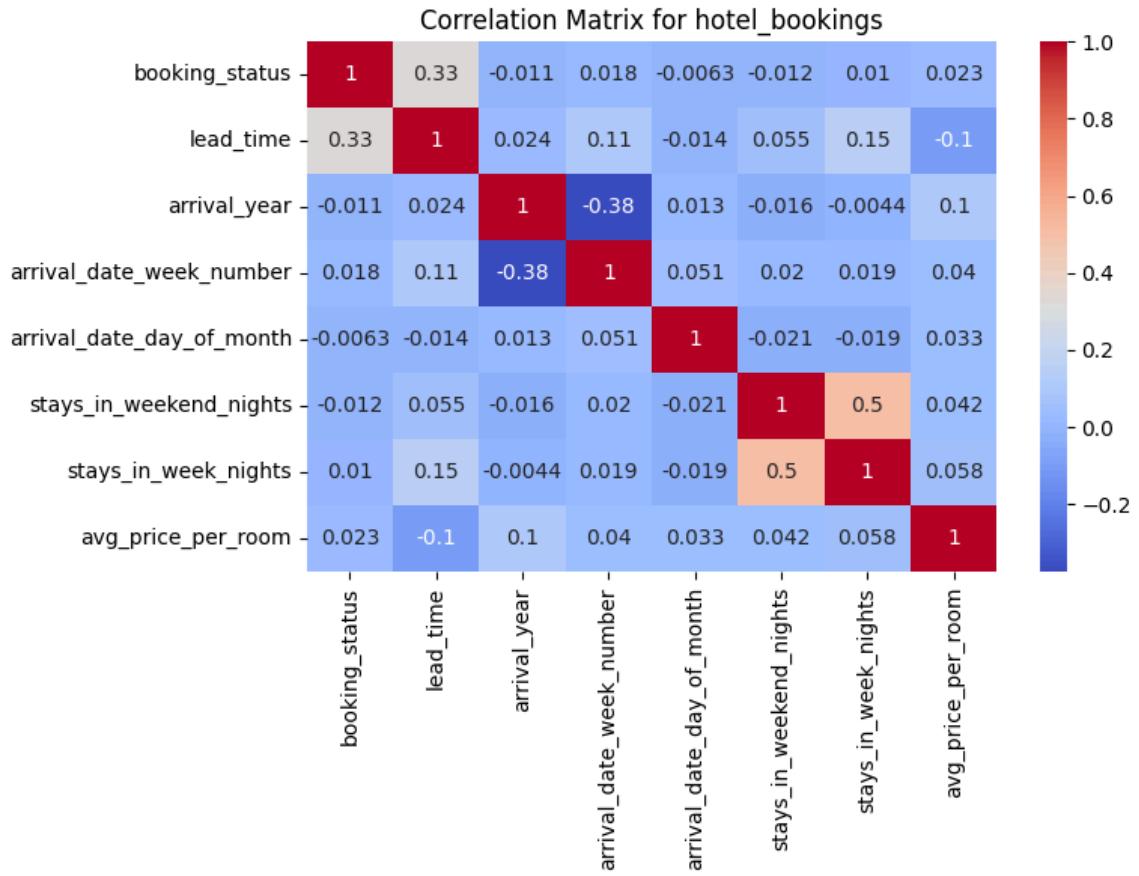


Figure 2: Correlation for Customer Reservation

Why:

This steps helps us in finding strong and weak relationships between continuous variables.

Insights:

- **Customer Reservations**
 - `stays_in_week_nights` and `stays_in_weekend_nights` show a **moderate positive correlation (0.18)** which seems reasonable as customers might extend their stays for the weekend if booked initially for the week days.
 - `lead_time` has **very weak or no correlation** with `avg_price_per_room` (≈ -0.06). This insight was kind of surprising as booking earlier had little or no effect in average room price.
 - Other relationships are near zero, confirming that most numeric fields (dates, lead time, price) are largely independent.
- **Hotel Bookings**
 - `stays_in_week_nights` and `stays_in_weekend_nights` have a **strong positive correlation (~0.50)** with same reasoning as above.
 - `lead_time` and `booking_status` are **positively correlated (~-0.33)** which means bookings made far in advance are slightly more likely to be canceled.
 - `lead_time` and `avg_price_per_room` show a **weak negative correlation (-0.10)** — earlier bookings tend to be marginally cheaper.
 - Other variables show near-zero correlations, suggesting minimal temporal dependencies.

Summary

Both the datasets are well structured and mostly clean. With EDA steps performed, we were able to find some useful insights from the datasets. These findings are useful in next steps of the data pre processing.

Data Merge & Integration

1. Dataset Overview

1.1 Source Datasets

We merged two hotel booking datasets: - Customer reservations: 36,275 records from customer booking system (no duplicates) - Hotel bookings: 78,703 records from hotel management system (1 duplicate removed during cleaning → 78,702 records) - Total merged: 114,977 records

2. Key Challenges

2.1 Different Column Names for Same Data

Booking_ID	N/A	Hotel data lacks booking identifiers
arrival_date	arrival_date_day_of_month	Different naming for same field

2.2 Different Data Formats

booking_status	Text ('Canceled', 'Not_Canceled')	Binary (0, 1)
arrival_month	Numeric (1-12)	Text ('January', 'February', ...)

2.3 Missing Columns

Customer Reservations	hotel, country, email, arrival_date_week_number	Added (hotel, country, email as NULL; week_number calculated)
Hotel Bookings	booking_id	Generated with INN prefix and offset

3. Data Cleaning Decisions

3.1 Customer Reservations Dataset

Column Standardization

- Renamed `Booking_ID` → `booking_id` for consistency
- Renamed `arrival_date` → `arrival_date_day_of_month` for clarity

Missing Column Handling

hotel	NULL	Customer Reservations system doesn't track hotel type
country	NULL	Geographic information not available
email	NULL	Privacy/data availability limitation
arrival_date_week_number	Computed calculated	from arrival_year, arrival_month, arrival_date_day_of_month

Week Number Calculation

The customer reservations dataset doesn't include week numbers, but we can derive them from the existing date components:

```
# Construct date from year, month, day components
date_string = concat_ws('-', year, padded_month, padded_day)
arrival_date = to_date(date_string)

# Extract ISO week number (1-53)
arrival_date_week_number = weekofyear(arrival_date)
```

Why:

- Maximizes data completeness by deriving missing information
- Uses ISO 8601 week numbering standard
- Enables consistent week-based analysis across both datasets
- Using NULL for hotel, country, and email follows database best practices for truly missing data

Duplicate Removal

- Checked for duplicate rows using all fields
- No duplicates found in customer reservations dataset

3.2 Hotel Bookings Dataset

Booking Status Standardization

- Original:** Binary (0 = not canceled, 1 = canceled)
- Target:** Text ('Not_Canceled', 'Canceled')

```
when(col('booking_status') == 0, 'Not_Canceled').otherwise('Canceled')
```

Why:

- Matches customer reservations dataset format

- Self-documenting (no need to remember 0/1 mapping)
 - Prevents accidental numeric operations on categorical data
-

Arrival Month Conversion

- **Original:** Text ('January', 'February', ..., 'December')
- **Target:** Numeric (1-12)

Why:

- Matches customer reservations dataset format
 - Enables numeric operations (sorting, filtering, calculations)
 - Reduces storage space
 - Facilitates time-series analysis
-

Booking ID Generation

- **Format:** INN50000, INN50001, INN50002, ...
- **Offset:** Starting at 50,000

Why:

Customer Reservations Dataset Booking IDs range from INN00001 to INN36275. Starting hotel IDs at 50,000 prevents ID collision while maintaining the same INN prefix pattern.

Duplicate Removal

- Found 1 duplicate record in Hotel Bookings Dataset
 - Removed during cleaning process
-

3.3 Final Merged Dataset Transformations

Boolean Conversion

- **Field:** booking_status → is_canceled
- **Format:** Boolean (True/False)

Why:

- Database-ready format (PostgreSQL BOOLEAN type)
- More efficient storage (1 byte vs 12+ bytes for text)
- Enables boolean operators in queries
- Clearer semantic meaning

ID Simplification

- **Field:** booking_id → id

Why:

Shorter, simpler column name for the primary identifier.

4. Schema Alignment Strategy

4.1 Unified Column Set

The following 14 columns were standardized across both datasets:

booking_id	String	Unique booking identifier
hotel	String	Hotel type (or NULL)
booking_status	String	Cancellation status
lead_time	Integer	Days between booking and arrival
arrival_year	Integer	Year of arrival
arrival_month	Integer	Month of arrival (1-12)
arrival_date_week_number	Integer	ISO week number (or calculated)
arrival_date_day_of_month	Integer	Day of month
stays_in_weekend_nights	Integer	Weekend nights booked
stays_in_week_nights	Integer	Weekday nights booked
market_segment_type	String	Market segment
country	String	Country code (or NULL)
avg_price_per_room	Double	Average room price
email	String	Customer email (or NULL)

5. Merge Methodology

5.1 Union vs Join Decision

Chosen Approach: UNION (vertical concatenation)

```
merged_df = customer_aligned.union(hotel_aligned)
```

Why:

- Datasets represent independent booking sources
- No natural join key exists between them
- Goal is to combine all records, not match them
- Preserves all records from both sources

Alternative Considered: JOIN operation

Why Rejected: Would require matching keys and could lose records; datasets are complementary, not overlapping.

5.2 Duplicate Detection and Removal

Duplicates are removed at the individual dataset level before merging: - **Customer Reservations:** Checked and found 0 duplicates - **Hotel Bookings:** Found and removed 1 duplicate record (Lisa_M@gmail.com with identical booking details)

6. Final Schema

6.1 Merged Dataset Structure

Total Columns: 14

<code>id</code>	String	Primary identifier
<code>hotel</code>	String	Nullable
<code>is_canceled</code>	Boolean	Not null
<code>lead_time</code>	Integer	Not null
<code>arrival_year</code>	Integer	Not null
<code>arrival_month</code>	Integer	Not null
<code>arrival_date_week_number</code>	Integer	Not null
<code>arrival_date_day_of_month</code>	Integer	Not null
<code>stays_in_weekend_nights</code>	Integer	Not null
<code>stays_in_week_nights</code>	Integer	Not null
<code>market_segment_type</code>	String	Not null
<code>country</code>	String	Nullable
<code>avg_price_per_room</code>	Double	Not null
<code>email</code>	String	Nullable

7. Key Decisions Summary

7.1 Data Quality

- **NULL Representation:** Missing data stored as NULL (blank in CSV) rather than placeholder strings
- **Duplicate Handling:** Business-logic-based deduplication catching semantic duplicates
- **ID Strategy:** Offset-based ID generation preventing collisions between datasets

7.2 Database Readiness

- **Boolean Fields:** booking_status converted to is_canceled boolean
- **Consistent Types:** All data types aligned and validated
- **Primary Key:** Unique id field for every record

7.3 Data Lineage

- **Customer Reservation Records:** INN00001 - INN36275 (36,275 records)
- **Hotel Booking Records:** INN50000 - INN128702 (78,702 records after deduplication)
- **Total Records:** 114,977 unique bookings

8. Summary

The data merge process successfully unified two independent booking systems into a single, clean dataset suitable for analysis and database import. Key achievements include:

1. **Schema Harmonization:** Resolved naming and format inconsistencies
2. **Data Quality:** Removed duplicates and standardized missing data handling
3. **Database Compatibility:** Converted to database-friendly formats (booleans, NULLs)
4. **Data Integrity:** Prevented ID collisions through offset strategy
5. **Complete Traceability:** All decisions documented with clear rationale

The resulting merged dataset provides a comprehensive view of hotel bookings across both systems, ready for downstream analysis and database integration.